

A Randomized Scheduler with Probabilistic Guarantees of Finding Bugs

by

Sebastian Burckhardt, Pravesh Kothari, Madanlal Musuvathi, Santosh Nagarakatte

Multicore Acceleration of Priority-Based Schedulers for Concurrency Bug Detection

by

Santosh Nagarakatte, Sebastian Burckhardt, Milo M. K. Martin, Madanlal Musuvathi

Präsentation

Manuel Maier

Seminar Concurrency and Memory Management

Department of Computer Sciences


Seminarleiter: Univ.-Prof. Dr. Christoph Kirsch

Introduction

- What is the problem?
 - Testing concurrency phenomena
 - How to find buggy code paths/code schedules?
- State of the art
 - Stress testing
 - Heuristic-directed testing
 - Systematic scheduling
 - Randomized scheduling(e.g.: PCT)
- What is a randomized scheduler?
 - PCT provides guaranteed probability of $1/(nk^{d-1})$
 - With n number of threads, k number of instructions and d depth.

Exkurs (Test) Coverage

- What is „Test Coverage“?



- What is the problem with coverage concerning Concurrency Testing?
- Image taken from [3]

Controlled scheduling

- Schedule selection
 - Choose the one schedule that causes an error
- Controlled scheduling
 - Less common interleavings.
 - Os threads controlled in User Mode
 - A scheduling policy → Determine the thread that makes progress

Needlepoint

- To test efficacy of previously proposed concurrency bug detectors.
- Mechanisms
 - Instrumenting synchronization operations
 - Blocking information
 - Fairness and starvation freedom
- Policies

Previously proposed concurrency bug detectors

- Random sleep(RS)
- Preemption always(PA)
- Preemption bounding(PB)
- Atom fuzzer(AF)
- PCT

Previously proposed concurrency bug detectors(2)




Figure 1. Bug detection abilities on common concurrency bugs for five different scheduling policies described in Section 3.2: Random Sleep (RS), Preemption Always (PA), Preemption Bounding (PB), AtomFuzzer (AF), and Probabilistic Concurrency Testing (PCT).

- Image taken from [2]

Typical concurrency bugs

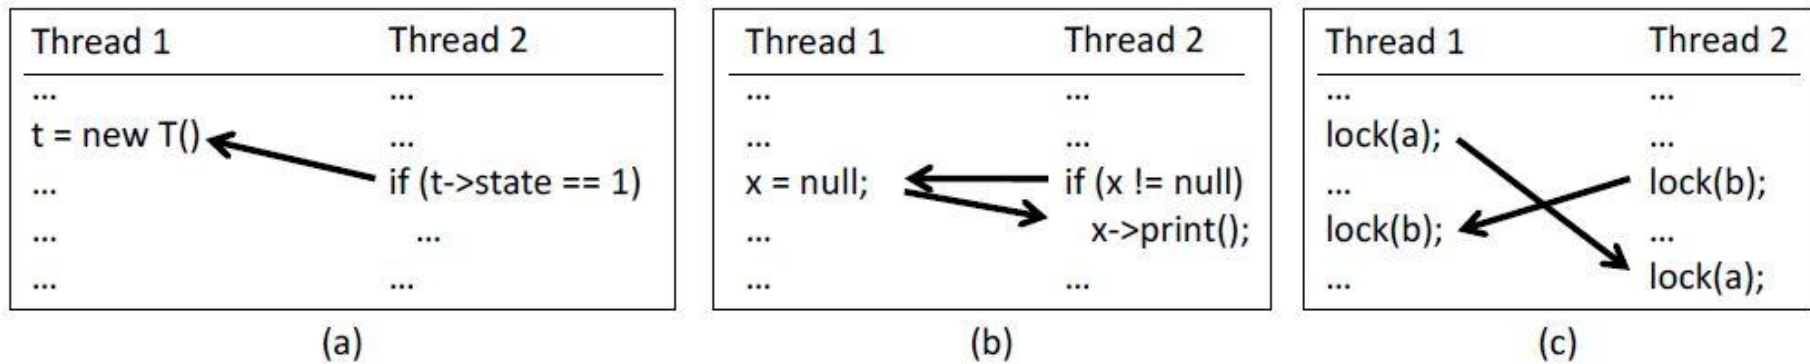


Figure 1. Three typical concurrency bugs, and ordering edges sufficient to find each. (A) This ordering bug manifests whenever the test by thread 2 is executed before the initialization by thread 1. (B) This atomicity violation manifests whenever the test by thread 2 is executed before the assignment by thread 1, and the latter is executed before the method call by thread 2. (C) This deadlock manifests whenever thread 1 locks a before thread 2, and thread 2 locks b before thread 1.

- Image taken from [1]

Bug Depth(d)

- Metric to classify concurrency bugs
- Minimum number of ordering constraints sufficient to find the bug.
 - Ordering bugs \rightarrow Depth 1
 - Atomicity violations \rightarrow Depth 2
 - Deadlocks(circular acquisition) \rightarrow Depth $n(n \text{ Threads})$
- Preemption bound($d - 1$):
Smallest number of preemptions to find bug




Figure 3. A variation of the example in Fig. 1(a), with the same bug depth of $d = 1$. Unlike in the other example, however, this bug requires Thread 1 to be preempted right after the instruction that sets the event e , and thus has a preemption bound of 1.

- Image taken from [1]

Bug Depth(2)

Program	Lines of code	Schedule points	Threads	Bug type	Bug depth
Pbzip2	15,188	1210	3	Ordering violation	2
Memc	11,182	845	4	Atomicity violation	2
t+15		2300			
t+25		3400			
t+35		3900			
am		79132			
WSQ-1	541	1916	4	Atomicity violation	3
WSQ-2		1086			2
WSQ-3		1717			2
Trans	33,622	38118	2	Ordering violation	1
NSPR	1,100	5361	3	Deadlock	2

Table 1. Concurrency bugs used for NeedlePoint's evaluation that we obtained from prior research [5, 10, 16, 27, 28]. WSQ is an implementation of the work stealing queue with lock free data structures. WSQ-1, WSQ-2, and WSQ-3 are the three distinct concurrency bugs in the WSQ implementation. Memc is the memcached daemon. Trans is the Transmission BitTorrent client.

- Image taken from [2]

Naive Randomization

- Flip coin each step




Figure 5. A program with two bugs of depth 1 that are hard to find with naive randomized schedulers that flip a coin in each step. PCT finds both these bugs with a probability $1/2$.

- Image taken from [1]

PCT –

Probabilistic Concurrency Testing

- Randomized scheduler with guaranteed probability $(1/(nk^{d-1}))$
- Algorithm:
 1. Assign n priority values $d, d+1, \dots, d+n$ to the n threads randomly.
 2. Pick $d-1$ random priority change points k_1, \dots, k_{d-1} . Each k_i has associated priority value i .
 3. Schedule threads honoring their priorities. When thread reaches i -th change point, change priority of that thread to i .
- Drawback

Typical concurrency bugs and PCT




Figure 6. Illustration on how our randomized scheduler finds bugs of depth d , using the examples from Fig. 1. The scheduler assigns random initial thread priorities $\{d, \dots, d + n - 1\}$ (white circles) and randomly places $d - 1$ priority change points of values $\{1, \dots, d - 1\}$ (black circles) into the execution. The bug is found if the scheduler happens to make the random choices shown above.

- Image taken from [1]

PPCT

- Intuition
 - Two sets of priority threads: High \leftrightarrow Low
 - PPCT serializes execution of lower priority set threads.
 - Number of „Low“ threads is bounded by d .
- Coverage guarantees
- Starvation and Priority Based Scheduling

PPCT Algorithm(1)

- Pick a random low priority thread.
 - At the beginning of the program, pick a thread uniformly at random and assign it a priority d . In addition, insert this thread into the lower priority set L . Insert all other threads into the higher priority set H .

PPCT Algorithm(2)

- Pick random priority change points
 - At the beginning of the program, pick $d - 1$ priority change points k_1, \dots, k_{d-1} . Each k_i has an associated priority value of i

PPCT Algorithm(3)

- Scheduler choice
 - At each step, the scheduler picks any non-blocked thread in H to schedule. If H is empty or if all threads in H are blocked, the scheduler picks the highest priority thread in L .

PPCT Algorithm(4)

- Priority change
 - After each step, increment a step counter. If the step counter matches k_i for some i , change priority of the just executed thread to i and insert it into L .

PCT compared to PPCT – Bug Detection

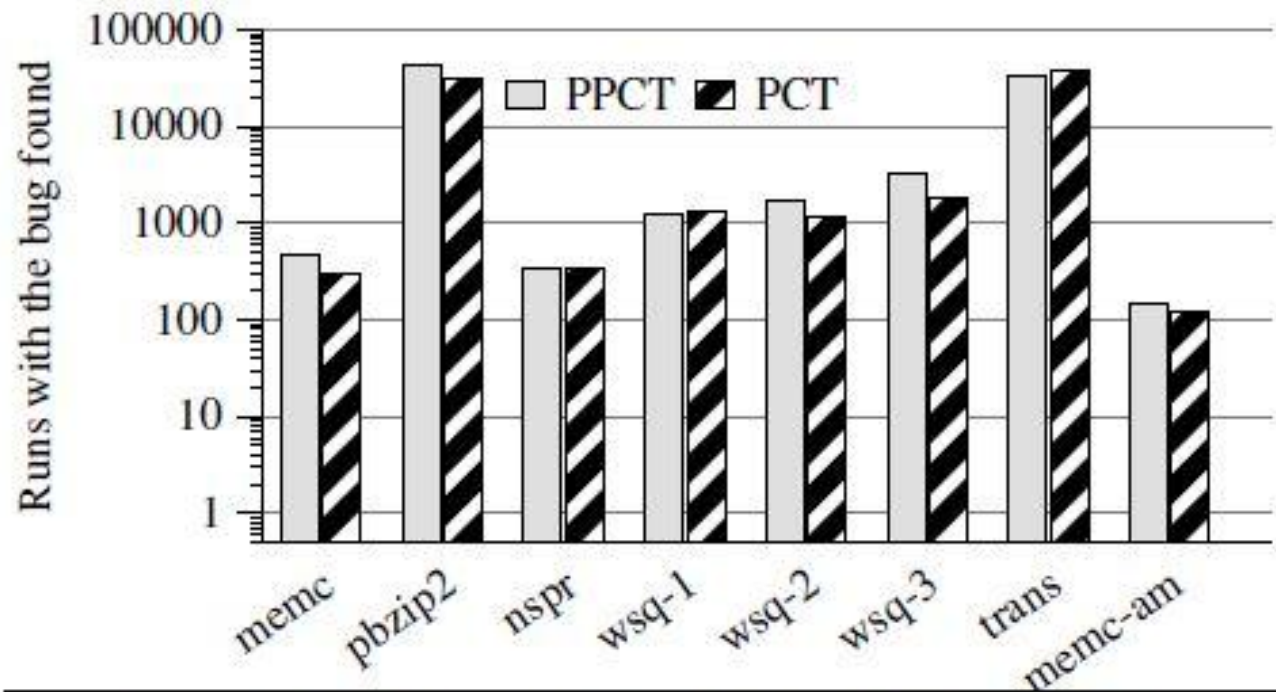


Figure 6. Bug detection of PPCT compared to PCT.

- Image taken from [2]

PCT compared to PPCT(2) - Slowdown




Figure 7. Runtime slowdown of PCT, PPCT on a single core and PPCT when compared to native multithreaded execution.

- Image taken from [2]

PCT compared to PPCT(3) - SpeedUp




Figure 8. PPCT speedup over PCT that on a multicore machine for explorations with various bug depths. Bug depth of each PPCT exploration is indicated at the bottom of the bar.

- Image taken from [2]

Referenzen

- **[1] A Randomized Scheduler with Probabilistic Guarantees of Finding Bugs**
by Sebastian Burckhardt, Pravesh Kothari, Madanlal Musuvathi, Santosh Nagarakatte
- **[2] Multicore Acceleration of Priority-Based Schedulers for Concurrency Bug Detection**
by Santosh Nagarakatte, Sebastian Burckhardt, Milo M. K. Martin, Madanlal Musuvathi
- **[3] Software Testing**
Udacity Course by John Regehr and Sean Bennett